

# Day 2 – 01 Explore Data

## Seminar plotting walk-through

We start the Day 2 session with the already prepared file `second_day_part2/data/dataset1_subset_long.csv`. This document simply shows how to open that file with base R commands and how to inspect what is inside before we move on to plotting. Run `second_day_part2/data/00_prepare_dataset.Rmd` first if you need to regenerate the CSV. Knit it from the repository root with:

```
R -e "rmarkdown::render('second_day_part2/scripts/01_explore_data.Rmd')"
```

### 1. Load the CSV

```
input_path <- file.path('..', 'data', 'dataset1_subset_long.csv')
long_df <- read.csv(input_path, stringsAsFactors = FALSE, check.names = FALSE)

cat('Rows:', nrow(long_df), '\nColumns:', ncol(long_df), '\n')

## Rows: 1536 \nColumns: 7 \n
```

### 2. Look at column names and a small preview

`colnames()` lists the headers exactly as they appear in the CSV and `head()` prints the first few rows. Encourage learners to read these outputs aloud so they begin to map column names to real-world meaning (genome name, SNP id, mouse ID, day, value, ...).

```
colnames(long_df)

## [1] "Genome"          "snp_id"           "Position"         "value"
## [5] "mouse_id"        "day"              "treatment_group"

head(long_df)

##                                     Genome      snp_id Position     value mouse_id day
## 1 Akkermansia_muciniphila_YL44 239840-C-G  239840 0.0000000   1683    0
## 2 Akkermansia_muciniphila_YL44 241793-A-G  241793 0.0495870   1683    0
## 3 Akkermansia_muciniphila_YL44 355328-A-T  355328 0.1381820   1683    0
## 4 Akkermansia_muciniphila_YL44 356291-C-A  356291 0.0000000   1683    0
## 5 Akkermansia_muciniphila_YL44 2351445-C-T 2351445 0.0000000   1683    0
## 6 Bacteroides_caecimuris_I48 1601848-T-C  1601848 0.0416090   1683    0
##   treatment_group
## 1 Control
## 2 Control
## 3 Control
## 4 Control
## 5 Control
## 6 Control
```

It can also help to peek at the SNP identifiers themselves to remind everyone what a typical label looks like:

```
head(unique(long_df$snp_id), n = 5)
## [1] "239840-C-G"  "241793-A-G"  "355328-A-T"  "356291-C-A"  "2351445-C-T"
```

### 3. Which genomes are present?

We often start by asking “Which genomes are represented here?”. `unique()` returns all distinct values, while `table()` counts how many rows belong to each genome.

```
unique(long_df$Genome)
## [1] "Akkermansia_muciniphila_YL44" "Bacteroides_caecimuris_I48"
table(long_df$Genome)

##
## Akkermansia_muciniphila_YL44      Bacteroides_caecimuris_I48
##                               320                      1216
```

Explain to students that each row contains the `Genome`, the unique SNP id, the mouse ID, the day of sampling, and the measured value.

### 4. How many SNPs per genome?

Counting the unique SNP identifiers per genome shows how much information we keep for each organism. The helper `tapply(values, groups, FUN)` applies a function to each subset of `values` defined by `groups`. Here we give it the vector of SNP ids and ask it to count how many unique ids (`length(unique(x))`) exist inside each genome group.

```
snp_per_genome <- tapply(long_df$snp_id, long_df$Genome, function(x) length(unique(x)))
snp_per_genome
## Akkermansia_muciniphila_YL44      Bacteroides_caecimuris_I48
##                               5                      19
```

### 6. How many measurements per mouse and per day?

```
table(long_df$mouse_id)

##
## 1683 1688 1692 1699
## 384 384 384 384
table(long_df$day)

##
## 0 4 9 14 18 23 30 37 44 49 53 58 63 67 72 79
## 96 96 96 96 96 96 96 96 96 96 96 96 96 96 96 96
```

For a combined view you can also build a two-way table. Again we use `table()`, this time with `mouse_id` on rows and `day` on columns. Reading across a row shows how many time points we have for a given mouse.

```
table(long_df$mouse_id, long_df$day)

##
##          0 4 9 14 18 23 30 37 44 49 53 58 63 67 72 79
## 1683 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24
## 1688 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24
```

```
##   1692 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24
##   1699 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24
```

Pick any cell and inspect those rows directly. For example, mouse 1683 on day 0 has the following measurements (all SNPs for that mouse/day):

```
subset_example <- long_df[long_df$mouse_id == '1683' & long_df$day == 0, ]
subset_example
```

	Genome	snp_id	Position	value	mouse_id	day
## 1	Akkermansia_muciniphila_YL44	239840-C-G	239840	0.000000	1683	0
## 2	Akkermansia_muciniphila_YL44	241793-A-G	241793	0.049587	1683	0
## 3	Akkermansia_muciniphila_YL44	355328-A-T	355328	0.138182	1683	0
## 4	Akkermansia_muciniphila_YL44	356291-C-A	356291	0.000000	1683	0
## 5	Akkermansia_muciniphila_YL44	2351445-C-T	2351445	0.000000	1683	0
## 6	Bacteroides_caecimuris_I48	1601848-T-C	1601848	0.041609	1683	0
## 7	Bacteroides_caecimuris_I48	1602949-A-G	1602949	0.000000	1683	0
## 8	Bacteroides_caecimuris_I48	2105155-C-T	2105155	0.113208	1683	0
## 9	Bacteroides_caecimuris_I48	2562529-T-C	2562529	0.000000	1683	0
## 10	Bacteroides_caecimuris_I48	2958703-C-A	2958703	0.000000	1683	0
## 11	Bacteroides_caecimuris_I48	2963503-C-A	2963503	0.000000	1683	0
## 12	Bacteroides_caecimuris_I48	2964459-G-T	2964459	0.000000	1683	0
## 13	Bacteroides_caecimuris_I48	2973657-A-G	2973657	0.000000	1683	0
## 14	Bacteroides_caecimuris_I48	2975719-A-G	2975719	0.000000	1683	0
## 15	Bacteroides_caecimuris_I48	2981479-T-G	2981479	0.000000	1683	0
## 16	Bacteroides_caecimuris_I48	2996527-C-A	2996527	0.000000	1683	0
## 17	Bacteroides_caecimuris_I48	2996770-A-G	2996770	0.000000	1683	0
## 18	Bacteroides_caecimuris_I48	3004527-A-T	3004527	0.000000	1683	0
## 19	Bacteroides_caecimuris_I48	3004659-G-A	3004659	0.000000	1683	0
## 20	Bacteroides_caecimuris_I48	3004706-A-G	3004706	0.000000	1683	0
## 21	Bacteroides_caecimuris_I48	3005011-C-A	3005011	0.000000	1683	0
## 22	Bacteroides_caecimuris_I48	3009589-C-T	3009589	0.000000	1683	0
## 23	Bacteroides_caecimuris_I48	3623124-G-A	3623124	0.173591	1683	0
## 24	Bacteroides_caecimuris_I48	4472955-A-G	4472955	0.000000	1683	0
##	treatment_group					
## 1	Control					
## 2	Control					
## 3	Control					
## 4	Control					
## 5	Control					
## 6	Control					
## 7	Control					
## 8	Control					
## 9	Control					
## 10	Control					
## 11	Control					
## 12	Control					
## 13	Control					
## 14	Control					
## 15	Control					
## 16	Control					
## 17	Control					
## 18	Control					
## 19	Control					
## 20	Control					

```

## 21      Control
## 22      Control
## 23      Control
## 24      Control

```

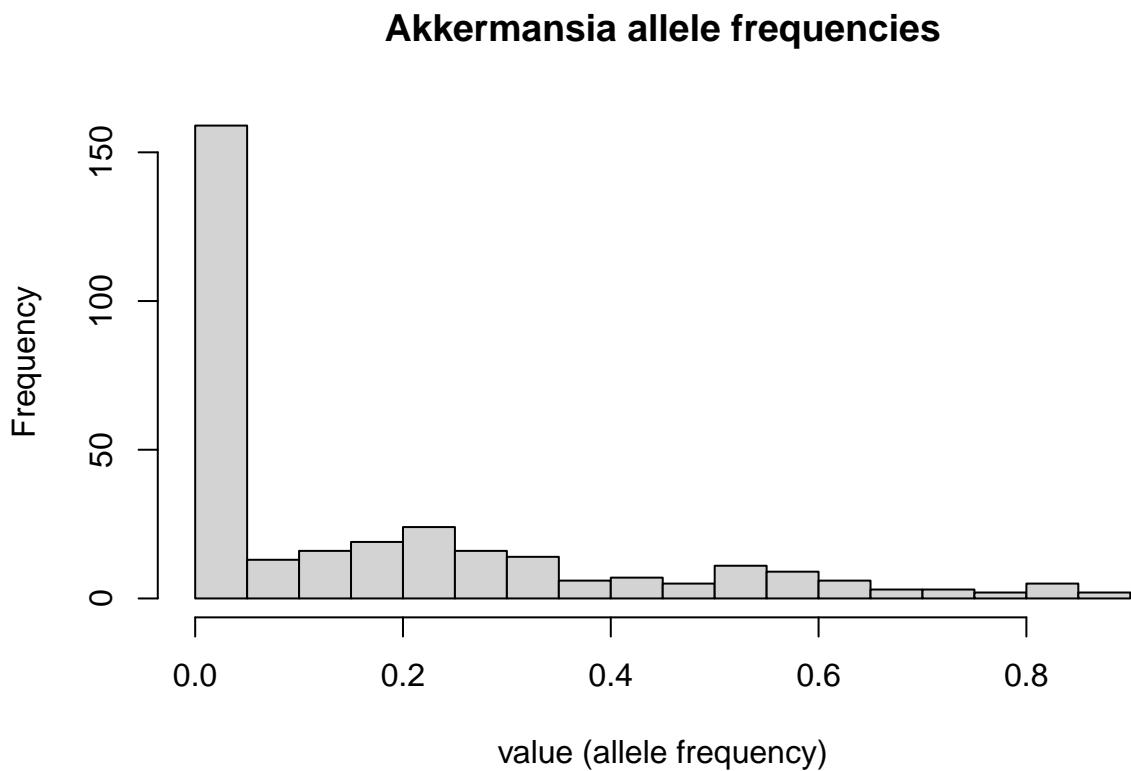
## 6. What do the allele-frequency values look like?

Start by focusing on a single genome (Akkermansia) so the histogram is easy to interpret, then compare it against the combined dataset.

```

akk_values <- long_df$value[long_df$Genome == 'Akkermansia_muciniphila_YL44']
hist(akk_values, breaks = 20, main = 'Akkermansia allele frequencies', xlab = 'value (allele frequency)')

```



```
summary(akk_values)
```

```

##   Min. 1st Qu. Median    Mean 3rd Qu.    Max.
## 0.00000 0.00000 0.05231 0.17419 0.27145 0.87125

```

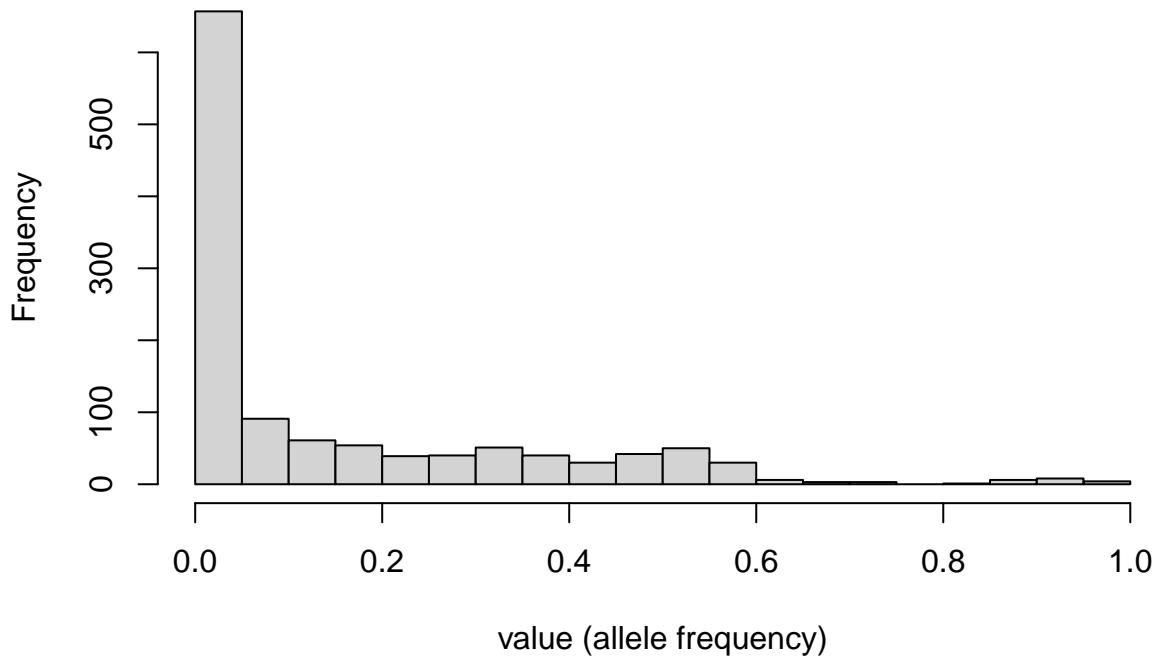
Do the same for Bacteroides (the other genome in this dataset).

```

bacto_values <- long_df$value[long_df$Genome == 'Bacteroides_caecimuris_I48']
hist(bacto_values, breaks = 20, main = 'Bacteroides allele frequencies', xlab = 'value (allele frequency)')

```

## Bacteroides allele frequencies



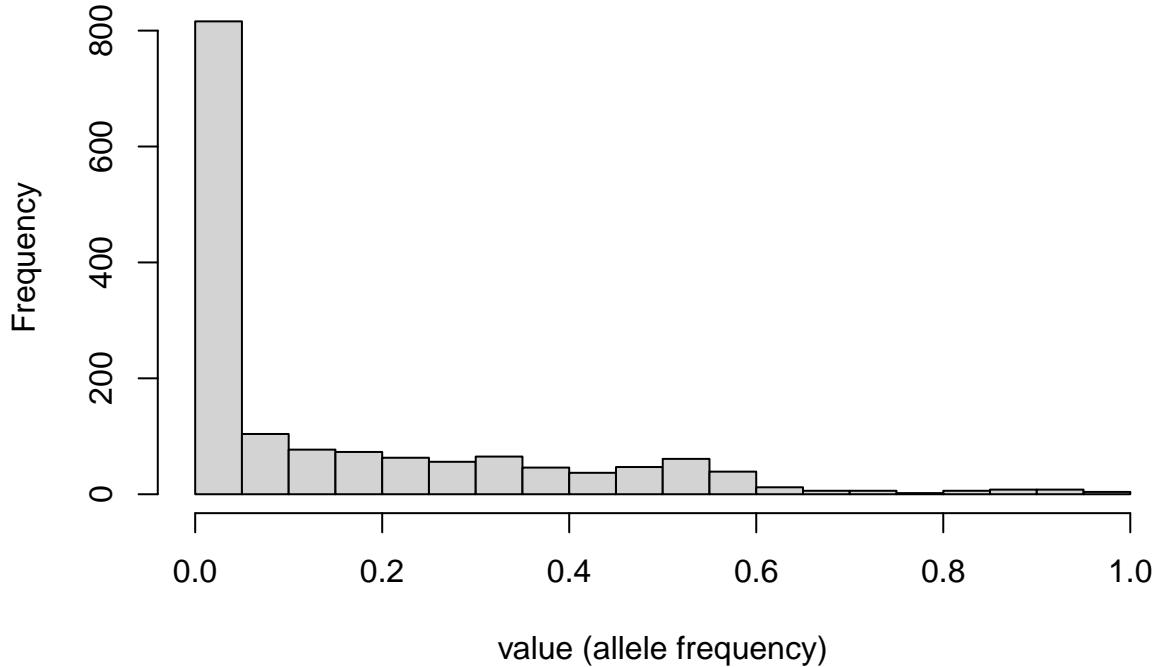
```
summary(bacto_values)
```

```
##      Min. 1st Qu. Median      Mean 3rd Qu.      Max.
## 0.00000 0.00000 0.03383 0.14854 0.26792 0.99122
```

Now look at the complete long table to see how the mixture of genomes changes the distribution.

```
hist(long_df$value, breaks = 30, main = 'All genomes: allele-frequency histogram', xlab = 'value (allele frequency)')
```

## All genomes: allele-frequency histogram



```
summary(long_df$value)
```

```
##      Min. 1st Qu. Median    Mean 3rd Qu.    Max.
## 0.00000 0.00000 0.03941 0.15388 0.27070 0.99122
```

To compare genomes explicitly, break the summary down by genome using `aggregate(value ~ Genome, ...)`. Read `value ~ Genome` as “value grouped by Genome”. The column on the left must be numeric because we are calling `summary()` on it.

```
value_by_genome <- aggregate(value ~ Genome, data = long_df, function(x) summary(x))
value_by_genome
```

```
##                               Genome value.Min. value.1st Qu. value.Median value.Mean
## 1 Akkermansia_muciniphila_YL44 0.00000000 0.00000000 0.0523090 0.1741893
## 2 Bacteroides_caecimuris_I48  0.00000000 0.00000000 0.0338305 0.1485411
##   value.3rd Qu. value.Max.
## 1     0.2714463 0.8712450
## 2     0.2679232 0.9912180
```

## 7. How many missing values?

Count how many NA values appear in `value` and, if any, locate them by mouse/day.

```
na_total <- sum(is.na(long_df$value))
na_total
```

```
## [1] 0
na_by_mouse_day <- with(long_df, tapply(value, list(mouse_id, day), function(x) sum(is.na(x))))
na_by_mouse_day
```

```
##      0 4 9 14 18 23 30 37 44 49 53 58 63 67 72 79
## 1683 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```

```
## 1688 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
## 1692 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
## 1699 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```

## 8. Treatment-group overview

The long table now contains a `treatment_group` column. Start by listing the distinct labels and counting how many rows belong to each group.

```
unique(long_df$treatment_group)  
  
## [1] "Control"      "Ciprofloxacin" "Tetracyclin"   "Vancomycin"  
  
table(long_df$treatment_group)  
  
##  
## Ciprofloxacin      Control    Tetracyclin    Vancomycin  
##            384           384           384           384
```

Which treatment groups appear for each mouse? Use a cross-tabulation to reinforce how mice were assigned.

```
table(long_df$mouse_id, long_df$treatment_group)
```

```
##  
##      Ciprofloxacin Control Tetracyclin Vancomycin  
## 1683          0     384          0          0  
## 1688          384          0          0          0  
## 1692          0          0     384          0  
## 1699          0          0          0     384
```

You can also inspect specific days. For example, day 30 only:

```
subset_day30 <- long_df[long_df$day == 30, c('mouse_id', 'treatment_group')]  
unique(subset_day30)  
  
##      mouse_id treatment_group  
## 97      1683        Control  
## 481      1688    Ciprofloxacin  
## 865      1692      Tetracyclin  
## 1249     1699      Vancomycin
```

## 9. Focus on a single genome (Akkermansia)

Subsetting is a great way to answer specific questions. Below we keep only `Akkermansia_muciniphila_YL44` (using a logical comparison inside the square brackets) and inspect the resulting table. This demonstrates how to focus on one organism without altering the original data frame.

Explain that `long_df$Genome == 'Akkermansia_muciniphila_YL44'` evaluates to a vector of TRUE/FALSE values—TRUE where the genome matches that text and FALSE everywhere else. When we place that logical vector inside `[...]`, R keeps only the TRUE rows.

```
akk_df <- long_df[long_df$Genome == 'Akkermansia_muciniphila_YL44', ]  
cat('Rows for Akkermansia:', nrow(akk_df), '\n')  
  
## Rows for Akkermansia: 320
```

```
head(akk_df)  
  
##      Genome      snp_id Position      value mouse_id day  
## 1 Akkermansia_muciniphila_YL44 239840-C-G 239840 0.000000      1683  0
```

```

## 2 Akkermansia_muciniphila_YL44 241793-A-G 241793 0.049587 1683 0
## 3 Akkermansia_muciniphila_YL44 355328-A-T 355328 0.138182 1683 0
## 4 Akkermansia_muciniphila_YL44 356291-C-A 356291 0.000000 1683 0
## 5 Akkermansia_muciniphila_YL44 2351445-C-T 2351445 0.000000 1683 0
## 25 Akkermansia_muciniphila_YL44 239840-C-G 239840 0.000000 1683 14
##   treatment_group
## 1       Control
## 2       Control
## 3       Control
## 4       Control
## 5       Control
## 25      Control

# After subsetting we can still use `table()` to see how measurements are
# distributed for this genome alone.





```

```
## [1] TRUE TRUE TRUE TRUE TRUE FALSE
```

Reading this out loud (“TRUE, TRUE, FALSE, …”) reinforces that the comparison creates a logical mask; only the TRUE rows survive when we subset with [...].

Having these quick checks documented makes it easy to reassure learners that we understand the input file before moving into the plotting notebook (`second_day_part2/scripts/02_simple_heatmap.Rmd`).