

Day 2 – 01 Explore Data

Seminar plotting walk-through

We start the Day 2 session with the already prepared file `second_day_part2/data/dataset1_subset_long.csv`. This document simply shows how to open that file with base R commands and how to inspect what is inside before we move on to plotting. Run `second_day_part2/data/00_prepare_dataset.Rmd` first if you need to regenerate the CSV. Knit it from the repository root with:

```
R -e "rmarkdown::render('second_day_part2/scripts/01_explore_data.Rmd')"
```

1. Load the CSV

```
input_path <- file.path('..', 'data', 'dataset1_subset_long.csv')
long_df <- read.csv(input_path, stringsAsFactors = FALSE, check.names = FALSE)

cat('Rows:', nrow(long_df), '\nColumns:', ncol(long_df), '\n')

## Rows: 1536 \nColumns: 7 \n
```

2. Look at column names and a small preview

`colnames()` lists the headers exactly as they appear in the CSV and `head()` prints the first few rows. Encourage learners to read these outputs aloud so they begin to map column names to real-world meaning (genome name, SNP id, mouse ID, day, value, ...).

```
colnames(long_df)

## [1] "Genome"          "snp_id"           "Position"         "value"
## [5] "mouse_id"        "day"              "treatment_group"

head(long_df)

##                                     Genome      snp_id Position     value mouse_id day
## 1 Akkermansia_muciniphila_YL44 239840-C-G  239840 0.0000000   1683    0
## 2 Akkermansia_muciniphila_YL44 241793-A-G  241793 0.0495870   1683    0
## 3 Akkermansia_muciniphila_YL44 355328-A-T  355328 0.1381820   1683    0
## 4 Akkermansia_muciniphila_YL44 356291-C-A  356291 0.0000000   1683    0
## 5 Akkermansia_muciniphila_YL44 2351445-C-T 2351445 0.0000000   1683    0
## 6 Bacteroides_caecimuris_I48 1601848-T-C  1601848 0.0416090   1683    0
##   treatment_group
## 1 Control
## 2 Control
## 3 Control
## 4 Control
## 5 Control
## 6 Control
```

It can also help to peek at the SNP identifiers themselves to remind everyone what a typical label looks like:

```
head(unique(long_df$snp_id), n = 5)
## [1] "239840-C-G"  "241793-A-G"  "355328-A-T"  "356291-C-A"  "2351445-C-T"
```

3. Which genomes are present?

We often start by asking “Which genomes are represented here?”. `unique()` returns all distinct values, while `table()` counts how many rows belong to each genome.

```
unique(long_df$Genome)
## [1] "Akkermansia_muciniphila_YL44" "Bacteroides_caecimuris_I48"
table(long_df$Genome)

##
## Akkermansia_muciniphila_YL44      Bacteroides_caecimuris_I48
##                               320                      1216
```

Explain to students that each row contains the `Genome`, the unique SNP id, the mouse ID, the day of sampling, and the measured value.

4. How many SNPs per genome?

Counting the unique SNP identifiers per genome shows how much information we keep for each organism. The helper `tapply(values, groups, FUN)` applies a function to each subset of `values` defined by `groups`. Here we give it the vector of SNP ids and ask it to count how many unique ids (`length(unique(x))`) exist inside each genome group.

```
snp_per_genome <- tapply(long_df$snp_id, long_df$Genome, function(x) length(unique(x)))
snp_per_genome
## Akkermansia_muciniphila_YL44      Bacteroides_caecimuris_I48
##                               5                      19
```

6. How many measurements per mouse and per day?

```
table(long_df$mouse_id)

##
## 1683 1688 1692 1699
## 384 384 384 384
table(long_df$day)

##
## 0 4 9 14 18 23 30 37 44 49 53 58 63 67 72 79
## 96 96 96 96 96 96 96 96 96 96 96 96 96 96 96 96
```

For a combined view you can also build a two-way table. Again we use `table()`, this time with `mouse_id` on rows and `day` on columns. Reading across a row shows how many time points we have for a given mouse.

```
table(long_df$mouse_id, long_df$day)

##
##          0 4 9 14 18 23 30 37 44 49 53 58 63 67 72 79
## 1683 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24
## 1688 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24
```

```
##   1692 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24
##   1699 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24 24
```

Pick any cell and inspect those rows directly. For example, mouse 1683 on day 0 has the following measurements (all SNPs for that mouse/day):

```
subset_example <- long_df[long_df$mouse_id == '1683' & long_df$day == 0, ]
subset_example
```

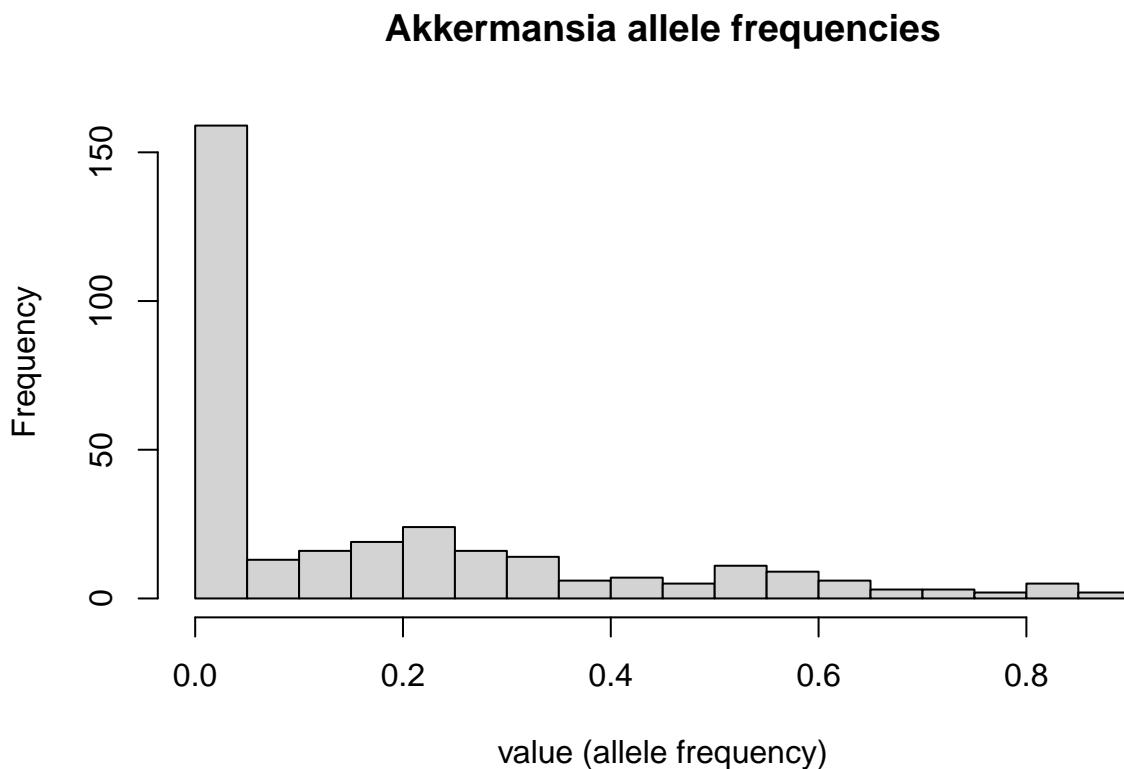
	Genome	snp_id	Position	value	mouse_id	day
## 1	Akkermansia_muciniphila_YL44	239840-C-G	239840	0.000000	1683	0
## 2	Akkermansia_muciniphila_YL44	241793-A-G	241793	0.049587	1683	0
## 3	Akkermansia_muciniphila_YL44	355328-A-T	355328	0.138182	1683	0
## 4	Akkermansia_muciniphila_YL44	356291-C-A	356291	0.000000	1683	0
## 5	Akkermansia_muciniphila_YL44	2351445-C-T	2351445	0.000000	1683	0
## 6	Bacteroides_caecimuris_I48	1601848-T-C	1601848	0.041609	1683	0
## 7	Bacteroides_caecimuris_I48	1602949-A-G	1602949	0.000000	1683	0
## 8	Bacteroides_caecimuris_I48	2105155-C-T	2105155	0.113208	1683	0
## 9	Bacteroides_caecimuris_I48	2562529-T-C	2562529	0.000000	1683	0
## 10	Bacteroides_caecimuris_I48	2958703-C-A	2958703	0.000000	1683	0
## 11	Bacteroides_caecimuris_I48	2963503-C-A	2963503	0.000000	1683	0
## 12	Bacteroides_caecimuris_I48	2964459-G-T	2964459	0.000000	1683	0
## 13	Bacteroides_caecimuris_I48	2973657-A-G	2973657	0.000000	1683	0
## 14	Bacteroides_caecimuris_I48	2975719-A-G	2975719	0.000000	1683	0
## 15	Bacteroides_caecimuris_I48	2981479-T-G	2981479	0.000000	1683	0
## 16	Bacteroides_caecimuris_I48	2996527-C-A	2996527	0.000000	1683	0
## 17	Bacteroides_caecimuris_I48	2996770-A-G	2996770	0.000000	1683	0
## 18	Bacteroides_caecimuris_I48	3004527-A-T	3004527	0.000000	1683	0
## 19	Bacteroides_caecimuris_I48	3004659-G-A	3004659	0.000000	1683	0
## 20	Bacteroides_caecimuris_I48	3004706-A-G	3004706	0.000000	1683	0
## 21	Bacteroides_caecimuris_I48	3005011-C-A	3005011	0.000000	1683	0
## 22	Bacteroides_caecimuris_I48	3009589-C-T	3009589	0.000000	1683	0
## 23	Bacteroides_caecimuris_I48	3623124-G-A	3623124	0.173591	1683	0
## 24	Bacteroides_caecimuris_I48	4472955-A-G	4472955	0.000000	1683	0
##	treatment_group					
## 1	Control					
## 2	Control					
## 3	Control					
## 4	Control					
## 5	Control					
## 6	Control					
## 7	Control					
## 8	Control					
## 9	Control					
## 10	Control					
## 11	Control					
## 12	Control					
## 13	Control					
## 14	Control					
## 15	Control					
## 16	Control					
## 17	Control					
## 18	Control					
## 19	Control					
## 20	Control					

```
## 21      Control
## 22      Control
## 23      Control
## 24      Control
```

6. What do the allele-frequency values look like?

Start by focusing on a single genome (Akkermansia) so the histogram is easy to interpret, then compare it against the combined dataset.

```
akk_values <- long_df$value[long_df$Genome == 'Akkermansia_muciniphila_YL44']
hist(akk_values, breaks = 20, main = 'Akkermansia allele frequencies', xlab = 'value (allele frequency)')
```



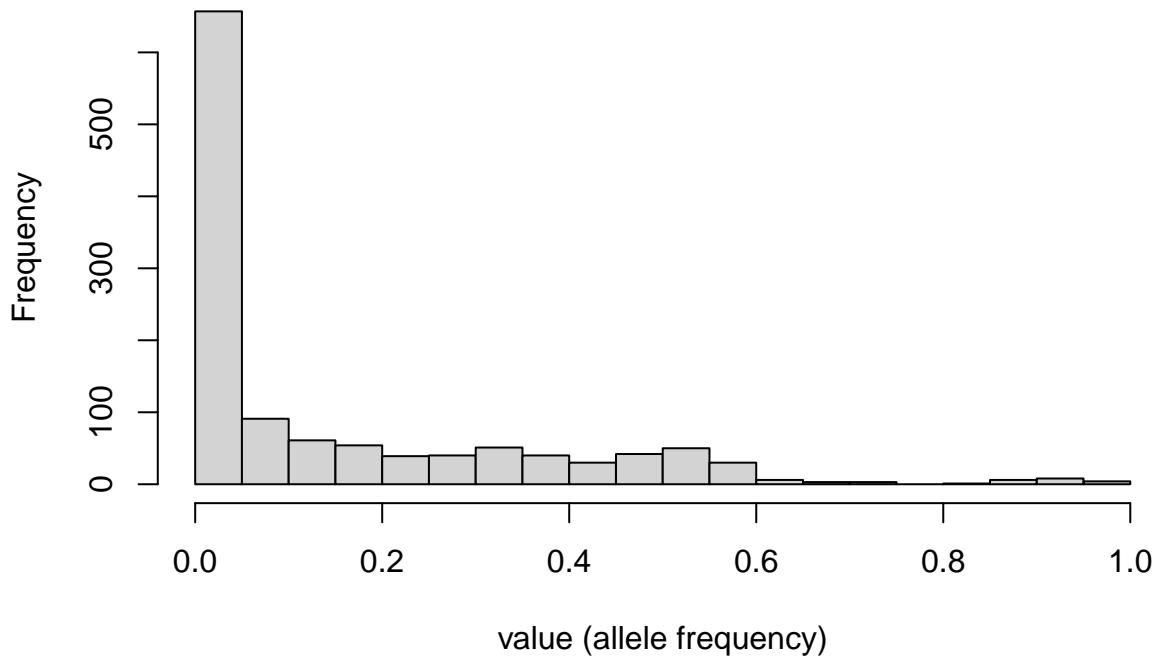
```
summary(akk_values)
```

```
##   Min. 1st Qu. Median    Mean 3rd Qu.    Max.
## 0.00000 0.00000 0.05231 0.17419 0.27145 0.87125
```

Do the same for Bacteroides (the other genome in this dataset).

```
bacto_values <- long_df$value[long_df$Genome == 'Bacteroides_caecimuris_I48']
hist(bacto_values, breaks = 20, main = 'Bacteroides allele frequencies', xlab = 'value (allele frequency)')
```

Bacteroides allele frequencies



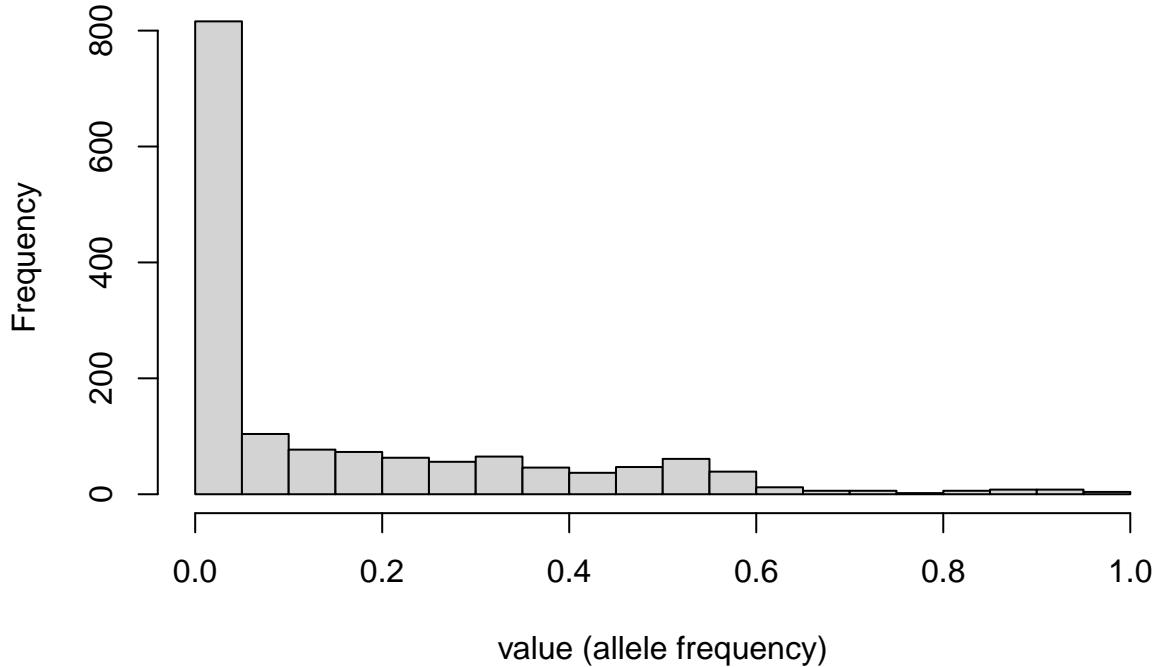
```
summary(bacto_values)
```

```
##      Min. 1st Qu. Median      Mean 3rd Qu.      Max.
## 0.00000 0.00000 0.03383 0.14854 0.26792 0.99122
```

Now look at the complete long table to see how the mixture of genomes changes the distribution.

```
hist(long_df$value, breaks = 30, main = 'All genomes: allele-frequency histogram', xlab = 'value (allele frequency)')
```

All genomes: allele-frequency histogram



```
summary(long_df$value)
```

```
##      Min. 1st Qu. Median    Mean 3rd Qu.    Max.
## 0.00000 0.00000 0.03941 0.15388 0.27070 0.99122
```

To compare genomes explicitly, break the summary down by genome using `aggregate(value ~ Genome, ...)`. Read `value ~ Genome` as “value grouped by Genome”. The column on the left must be numeric because we are calling `summary()` on it.

```
value_by_genome <- aggregate(value ~ Genome, data = long_df, function(x) summary(x))
value_by_genome
```

```
##                               Genome value.Min. value.1st Qu. value.Median value.Mean
## 1 Akkermansia_muciniphila_YL44 0.00000000 0.00000000 0.0523090 0.1741893
## 2 Bacteroides_caecimuris_I48  0.00000000 0.00000000 0.0338305 0.1485411
##   value.3rd Qu. value.Max.
## 1     0.2714463 0.8712450
## 2     0.2679232 0.9912180
```

7. How many missing values?

Count how many NA values appear in `value` and, if any, locate them by mouse/day.

```
na_total <- sum(is.na(long_df$value))
na_total
```

```
## [1] 0
na_by_mouse_day <- with(long_df, tapply(value, list(mouse_id, day), function(x) sum(is.na(x))))
na_by_mouse_day
```

```
##      0 4 9 14 18 23 30 37 44 49 53 58 63 67 72 79
## 1683 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```

```
## 1688 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
## 1692 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
## 1699 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```

8. Treatment-group overview

The long table now contains a `treatment_group` column. Start by listing the distinct labels and counting how many rows belong to each group.

```
unique(long_df$treatment_group)  
  
## [1] "Control"      "Ciprofloxacin" "Tetracyclin"   "Vancomycin"  
  
table(long_df$treatment_group)  
  
##  
## Ciprofloxacin      Control    Tetracyclin    Vancomycin  
##            384          384          384          384
```

Which treatment groups appear for each mouse? Use a cross-tabulation to reinforce how mice were assigned.

```
table(long_df$mouse_id, long_df$treatment_group)
```

```
##  
##      Ciprofloxacin Control Tetracyclin Vancomycin  
## 1683          0     384          0          0  
## 1688          384          0          0          0  
## 1692          0          0     384          0  
## 1699          0          0          0     384
```

You can also inspect specific days. For example, day 30 only:

```
subset_day30 <- long_df[long_df$day == 30, c('mouse_id', 'treatment_group')]  
unique(subset_day30)  
  
##      mouse_id treatment_group  
## 97      1683        Control  
## 481      1688    Ciprofloxacin  
## 865      1692      Tetracyclin  
## 1249     1699      Vancomycin
```

9. Focus on a single genome (Akkermansia)

Subsetting is a great way to answer specific questions. Below we keep only `Akkermansia_muciniphila_YL44` (using a logical comparison inside the square brackets) and inspect the resulting table. This demonstrates how to focus on one organism without altering the original data frame.

Explain that `long_df$Genome == 'Akkermansia_muciniphila_YL44'` evaluates to a vector of TRUE/FALSE values—TRUE where the genome matches that text and FALSE everywhere else. When we place that logical vector inside `[...]`, R keeps only the TRUE rows.

```
akk_df <- long_df[long_df$Genome == 'Akkermansia_muciniphila_YL44', ]  
cat('Rows for Akkermansia:', nrow(akk_df), '\n')  
  
## Rows for Akkermansia: 320  
  
head(akk_df)
```

```
##           Genome      snp_id Position      value mouse_id day  
## 1 Akkermansia_muciniphila_YL44 239840-C-G 239840 0.000000     1683  0
```

```

## 2 Akkermansia_muciniphila_YL44 241793-A-G 241793 0.049587 1683 0
## 3 Akkermansia_muciniphila_YL44 355328-A-T 355328 0.138182 1683 0
## 4 Akkermansia_muciniphila_YL44 356291-C-A 356291 0.000000 1683 0
## 5 Akkermansia_muciniphila_YL44 2351445-C-T 2351445 0.000000 1683 0
## 25 Akkermansia_muciniphila_YL44 239840-C-G 239840 0.000000 1683 14
##   treatment_group
## 1       Control
## 2       Control
## 3       Control
## 4       Control
## 5       Control
## 25      Control

# After subsetting we can still use `table()` to see how measurements are
# distributed for this genome alone.





```

```
## [1] TRUE TRUE TRUE TRUE TRUE FALSE
```

Reading this out loud (“TRUE, TRUE, FALSE, …”) reinforces that the comparison creates a logical mask; only the TRUE rows survive when we subset with [...].

Having these quick checks documented makes it easy to reassure learners that we understand the input file before moving into the plotting notebook (`second_day_part2/scripts/02_simple_heatmap.Rmd`).